

NVIDIA Corporation: Valuing the Engine of the AI Revolution

The Kimi K3 Shock — Case and Exhibits

Asaf Manela · Olin Business School, Washington University in St. Louis

Case written July 30, 2026

Valuation date: July 30, 2026 (market data as of the July 29, 2026 close)

Priya Raghunathan and Kestrel Ridge Partners are fictional. NVIDIA Corporation, Moonshot AI, and the other organizations named are real, and all financial and market data are drawn from public sources as of the dates shown. This case was prepared solely as a basis for class discussion. It is not a research report and not investment advice, and it should not be relied upon for any actual valuation of NVIDIA Corporation, its competitors, or the semiconductor and artificial intelligence industries. Some events are simplified or condensed for teaching purposes.

1. The Decision

NVIDIA closed on Wednesday, July 29, 2026 at **\$190.01**, down 3.6% on the day and down 20% from its 52-week high of \$236.54. Priya Raghunathan watched the print with the resignation of someone who had seen it coming and had not acted.

Raghunathan runs technology research at Kestrel Ridge Partners, a \$6.8 billion long/short equity fund. NVIDIA has been the fund's largest position for three years. It is now also its largest problem. Since the close on July 15 the stock has fallen 10.6%, erasing roughly \$545 billion of market value — more than the entire market capitalization of all but a handful of public companies.

The fund's investment committee meets Monday. The memo is due Friday.

"I don't want a view on the tape," the fund's CIO told her. "I want to know what this business is worth. Then tell me whether the market is paying more or less than that, and what has to be true for each answer to be right."

The immediate trigger for the selloff was not an earnings miss. NVIDIA's most recent quarter, reported May 20, 2026, was the strongest in its history: **\$81.6 billion of revenue, up 85% year over year**, with guidance for \$91 billion in the following quarter. The trigger was a piece of software released 8,000 miles away.

On **July 16, 2026**, Moonshot AI, a Beijing startup founded in 2023, released **Kimi K3** — a 2.8-trillion-parameter open-weight model that its own benchmarks, and several independent leaderboards, placed within striking distance of the best models from OpenAI and Anthropic. It was priced at \$15 per million output tokens — 30% of what the leading Western frontier model charges, and half of the next one down. On July 27 Moonshot published the full model weights on Hugging Face under its own Kimi K3 License — open enough for almost any enterprise to download and self-host, though large commercial redistributors face conditions.

Within two sessions the Philadelphia Semiconductor Index (SOX) had fallen roughly 10% — its worst week since April 2025 — and entered a bear market. Investors reached immediately for the analogy: this was DeepSeek all over again, the January 2025 episode in which a cheap Chinese model erased \$589 billion of NVIDIA's market capitalization in a single day.

Raghunathan is not convinced the analogy holds. DeepSeek's claim was that a frontier-class model could be *trained* for a small fraction of the expected cost — a direct attack on the premise that AI leadership requires enormous training clusters. Kimi K3's claim is different and, in some ways, stranger. K3 is not small. At 2.8 trillion parameters and roughly 1.56 terabytes of weights, it is the **largest** open-weight model ever released. It is more efficient per unit of capability, but it is enormous, and serving it at scale is a memory-bandwidth problem more than a raw-compute problem.

So which is it? Does a cheap, excellent, freely downloadable model reduce the world's demand for NVIDIA silicon, because less compute is needed per unit of intelligence? Or does it increase demand, because intelligence just got cheap enough for everyone to use constantly?

Jensen Huang, NVIDIA’s co-founder and chief executive, has been unambiguous. In public remarks after the release he called the Chinese models “excellent,” argued that “open-source models that are excellent should be used,” and summarized his position on the selloff in six words: “Free AI should be great for hardware.” A smarter, cheaper model attracts more users, and more users need more computers. On July 23 he went further, publishing an open letter — *Open Weights and American AI Leadership* — opposing proposals in Washington to restrict downloadable AI models.

Raghunathan’s job is to convert that argument into a number.

2. What Happened: A Timeline

Raghunathan’s first move whenever a position broke was always the same: build the timeline before building the model. Narratives assembled after the fact tend to explain whatever just happened. A timeline built from dated facts sometimes refuses to.

Date (2026)	Event
Jan 13–14	The Bureau of Industry and Security moves H200 exports to China from presumption of denial to case-by-case review; a presidential proclamation imposes a 25% Section 232 tariff on H200s imported into the United States, where each China-bound unit must be routed for inspection. Beijing responds within days, reportedly advising domestic technology firms against buying NVIDIA parts and instructing customs to block H200 entry.
Feb 25	NVIDIA reports fiscal 2026: revenue \$215.9B, up 65%; net income \$120.1B. Q1 FY27 guidance of \$78B.
Feb 26	NVIDIA confirms it has received licenses to ship a limited quantity of H200 accelerators to China. Whether any units will be allowed in remains unresolved.
Mar 16	At GTC, NVIDIA announces the Vera Rubin platform — seven co-designed chips, all in full production, with partner systems available in the second half of the year. Management raises cumulative Blackwell-plus-Rubin revenue visibility to roughly \$1 trillion through calendar 2027 , up from the \$500 billion figure disclosed in late 2025.
May 20	Q1 FY2027: revenue \$81.6B (+85% YoY), Data Center \$75.2B (+92%), non-GAAP EPS \$1.87, free cash flow \$48.6B. Dividend raised from \$0.01 to \$0.25 per quarter; \$80B added to buyback authorization. Q2 guidance \$91B, excluding any China data center compute revenue.
Jul 6–14	Semiconductor stocks weaken on valuation concerns after a roughly 130% twelve-month run in the SOX. NVDA closes at \$212.50 on July 15.
Jul 16	Moonshot AI unveils Kimi K3 , available the same day through its API and consumer app.
Jul 17	K3 is exhibited on the opening day of the World AI Conference in Shanghai, which runs July 17–20. Global semiconductor selloff; Asian chip indices fall sharply. NVDA closes at \$202.81 (–2.2%). The SOX ends the week down roughly 10% and in a bear market.
Jul 20–24	Partial rebound as Bank of America, UBS, and Morgan Stanley publish notes arguing the selloff is a buying opportunity. NVDA recovers to \$212.06 on July 22 before drifting to \$206.84 on July 24.
Jul 22	Michael Kratsios, Director of the White House Office of Science and Technology Policy, publicly accuses Moonshot of training K3 on NVIDIA GB300 systems hosted in Thailand, circumventing U.S. export controls. Moonshot is reported to be seeking additional Blackwell capacity for a successor model.
Jul 27	Moonshot publishes the full K3 weights on Hugging Face. Separately, reports emerge that NVIDIA is weighing a guarantee of as much as \$250 billion in OpenAI data center lease payments and up to \$350 billion in financing for OpenAI GPU purchases. “Circular financing” concerns resurface. NVDA closes at \$196.51 (–5.0%).
Jul 29	NVDA closes at \$190.01 , market capitalization \$4.60 trillion .

Laid out that way, the sequence bothered her. The market had already reversed the K3 shock by July 22 — the stock closed higher that day than it had the morning K3 launched. The damage came afterward, on the day the weights actually went up and the day the OpenAI financing reports landed. Whatever the market was repricing, it had taken more than a week to decide what it thought, which is not how a genuine surprise usually behaves.

3. What Kimi K3 Actually Is

Raghunathan had learned to distrust her own instinct to skip the technology. Twice before — with cloud margins in 2017, and with DeepSeek in 2025 — she had formed a view from summaries and been wrong in a way that a day of primary reading would have caught. So she spent Thursday on Moonshot’s technical report and the model card, trying to understand K3 at the level a generalist investor actually needs.

3.1 Architecture

Kimi K3 is a **mixture-of-experts (MoE)** model. A dense model of 2.8 trillion parameters would use all 2.8 trillion for every token it generates. An MoE model instead routes each token to a small subset of specialized sub-networks. K3 activates **16 of 896 experts** per token — roughly **104 billion active parameters** out of 2.8 trillion total.

This is the central economic fact about K3. The model must be **stored** in full — about 1.56 terabytes of weights in the four-bit format Moonshot trained it in — but only about 4% of it is **computed** on any given token.

Three other design choices matter:

- **Kimi Delta Attention (KDA)**, a hybrid linear attention mechanism. Moonshot’s earlier Kimi Linear research reported up to 6.3× faster decoding at million-token context lengths; the K3 release itself publishes no speedup figure.
- **Attention Residuals (AttnRes)**, which the underlying research reports reaches baseline quality at 1.25× the compute — roughly a 25% efficiency gain — for under 4% additional training overhead.
- **Quantization-aware training** from the supervised fine-tuning stage onward, using 4-bit weights and 8-bit activations. *Quantization* means storing each number in fewer bits, which shrinks the model in memory at some cost to precision. K3 was trained to be small in memory rather than compressed after the fact.

Moonshot claims roughly a **2.5× improvement in overall scaling efficiency** relative to its previous model, Kimi K2. The context window is 1,048,576 tokens — approximately a million tokens, or several thousand pages.

The sparsity was the part Raghunathan kept returning to. A model that stores 2.8 trillion parameters but computes 104 billion of them is cheap to *run* and expensive to *hold*. Those are different bills, paid to different suppliers.

3.2 Capability and Price

Dimension	Kimi K3	Leading Western frontier models
Output price (per million tokens)	\$15	\$30–\$50
Input price (per million tokens)	\$3.00 (\$0.30 cached)	Higher
Weights	Open, downloadable	Closed
Independent leaderboard standing	#1 Frontend Code Arena; #2 Vals AI; #3 Artificial Analysis Intelligence Index	Claude Fable 5 and GPT-5.6 Sol hold the top two places on most indices

Moonshot’s own technical documentation concedes that overall performance “still trails the most powerful proprietary models.” The competitive claim was not that K3 was best. It was that K3 was close enough, at 30–50% of the price depending on the comparator, and available to self-host.

That was a claim about the frontier labs’ pricing power, Raghunathan noted, not directly about NVIDIA’s. The two were connected, but not by a short chain.

3.3 Why “Open Weights” Is the Financially Interesting Part

A closed model is a service. You send tokens to a vendor and pay per token. An open-weight model is an **asset**. Any organization with sufficient hardware can download K3, run it inside its own firewall, fine-tune it on proprietary data, and pay nothing per token beyond the electricity and depreciation on the machines.

For a bank, an insurer, or an asset manager subject to data-residency rules, that distinction is not incremental — it is the difference between “we cannot use this” and “we can.” Kestrel Ridge’s own compliance team had rejected three vendors on precisely those grounds in the past year.

Which brought Raghunathan to the question she could not get around, and which she wrote at the top of a blank page: **when the marginal cost of intelligence falls, does the total amount spent on the hardware that produces intelligence rise or fall?**

3.4 The Memory Argument

Serving K3 at scale means holding roughly 1.56 TB of weights in fast memory, streaming them across a million-token context, and doing so for a workload dominated by *decoding* — generating output tokens one at a time — where the bottleneck is how fast weights can be moved, not how fast arithmetic can be performed. Eight AMD MI355X accelerators, each with 288 GB of high-bandwidth memory (HBM), provide 2.3 TB — enough to hold the weights plus working state. The equivalent NVIDIA configurations are similarly memory-defined.

In other words, K3's efficiency gains are concentrated in **compute per token**, while its resource demands are concentrated in **memory capacity and bandwidth**. If the industry's models trend toward K3's shape — enormous, sparse, long-context — then the binding constraint on AI infrastructure shifts from raw arithmetic throughput toward HBM, the advanced packaging that stacks it beside the processor, and the interconnect that links processors together.

That shift is visible in industry forecasts. The World Semiconductor Trade Statistics organization's Spring 2026 forecast projects the **memory** segment to grow roughly 250% year over year in 2026, exceeding \$800 billion, within a total semiconductor market of \$1.51 trillion. Bank of America's semiconductor team argued that K3 *reinforces* rather than undermines the memory thesis: cheap open models deployed inside enterprises increase demand for HBM, DRAM, and NAND.

NVIDIA sells memory-bandwidth-defined systems, not bare processors. Its Vera Rubin VR200 pairs 50 petaflops of FP4 compute with **22 TB/s of memory bandwidth** on a 288 GB footprint; the NVL72 rack configuration delivers 3.6 exaflops and 260 TB/s of all-to-all NVLink bandwidth.

Here Raghunathan found herself genuinely undecided, and she distrusted the confidence of everyone who was not. NVIDIA does not manufacture high-bandwidth memory; it buys it, which meant a memory-bound world raised its input costs and handed pricing power to Samsung, SK hynix, and Micron. But a 1.56-terabyte model has to be split across many accelerators, and the moment it is split, the *interconnect* becomes as binding as the memory — and interconnect was where NVIDIA's lead was widest. She could argue it either way, which was itself informative: the honest answer was that K3 pushed on NVIDIA's costs and its moat at the same time.

3.5 The Geopolitical Overlay

On July 22 the White House alleged that Moonshot trained K3 on NVIDIA GB300 systems located in Thailand — hardware Moonshot could not legally have imported into China, accessed remotely so that no chip crossed a Chinese border. In February 2026, Anthropic had separately accused Moonshot and two other Chinese labs of harvesting reasoning traces from its models through fraudulent accounts.

Both allegations, if accurate, cut in complicated directions for NVIDIA:

- They imply that **Chinese demand for NVIDIA hardware is real and unmet**, which is bullish for units if controls ever loosen.
- They imply **tighter enforcement is coming**, including possible restrictions on offshore remote access, which is bearish for a channel that is currently generating revenue.
- They raise the possibility of U.S. restrictions on downloadable model weights themselves — a policy Huang has publicly opposed, on the theory that a world running Chinese open models is a world that eventually buys Chinese hardware.

China contributes essentially none of NVIDIA's data center compute revenue today. The company excluded it entirely from Q2 FY2027 guidance. Sell-side estimates put the opportunity at \$4–10 billion annually if both governments permit it — meaningful, but under 3% of consensus FY2027 revenue.

Raghunathan set the China question aside deliberately. At that scale it was a call option, not a load-bearing assumption, and she had watched too many analysts spend their best hours on the most legible risk rather than the largest one.

4. NVIDIA Today

On Friday morning Raghunathan closed the browser tabs and went back to the filings. Whatever K3 meant, it would have to show up eventually in one of four places: revenue, margin, reinvestment, or the discount rate. Everything else was commentary.

4.1 The Business

NVIDIA designs accelerated computing systems and outsources manufacturing, principally to TSMC. As of fiscal 2026 it reports two market platforms:

Data Center — \$193.7 billion in fiscal 2026, 90% of revenue, up 68%. For the most recent quarter NVIDIA disclosed the split as: - *Compute* — \$60.4 billion in the quarter, up 77% year over year. GPUs and systems for training and inference. - *Networking* — \$14.8 billion, up **199%**. NVLink, InfiniBand, Spectrum Ethernet, and BlueField DPUs. This is the fastest-growing line in the company and the least appreciated part of the moat: at rack scale, the interconnect *is* the computer.

NVIDIA further divides Data Center demand into **Hyperscale** (roughly \$38 billion in Q1 FY2027, about half the segment) and **ACIE** — AI Clouds, Industrial, and Enterprise — at roughly \$37 billion, growing 31% sequentially. The diversification away from four or five hyperscale buyers toward sovereign programs, *neoclouds* — specialist GPU rental providers such as CoreWeave and Nebius — and enterprises is one of the most consequential changes in the business over the past two years.

Edge Computing — \$6.4 billion in Q1 FY2027, up 29%. Consolidates gaming GPUs, professional visualization, and automotive/robotics. Profitable and stable, and — Raghunathan concluded after twenty minutes she later judged wasted — immaterial to the valuation. At under 8% of revenue it could double or vanish without changing her answer.

4.2 The Numbers

Line item	FY2024	FY2025	FY2026	Q1 FY2027
Revenue	\$60.9B	\$130.5B	\$215.9B	\$81.6B
Growth	+126%	+114%	+65%	+85% YoY
GAAP gross margin	72.7%	75.0%	71.1%	74.9%
GAAP operating income	\$33.0B	\$81.5B	\$130.4B	\$53.5B
GAAP operating margin	54.1%	62.4%	60.4%	65.6%
GAAP net income	\$29.8B	\$72.9B	\$120.1B	\$58.3B
GAAP diluted EPS	\$1.19†	\$2.94	\$4.90	\$2.39
Operating cash flow	\$28.1B	\$64.1B	\$102.7B	\$50.3B
Free cash flow	\$27.0B*	\$60.7B	\$96.6B	\$48.6B

† Split-adjusted for the 10-for-1 split of June 2024. * NVIDIA defines free cash flow as operating cash flow less both purchases of, and principal payments on, property, equipment and intangible assets. The FY2024 figure is operating cash flow less capital expenditure only.

Three things in the fiscal 2026 numbers changed how Raghunathan thought about the position.

First, most of the gross margin dip was an artifact. Fiscal 2026's 71.1% GAAP gross margin looks like sharp compression against fiscal 2025's 75.0%. Q1 FY2026 carried a \$4.5 billion charge on H20 inventory stranded by export controls, which pushed that quarter to 60.5%; adding the charge back lifts the full year to **73.2%**. The exit rate tells the cleaner story: Q4 FY2026 was 75.0% and Q1 FY2027 was 74.9%. So there is no compression in the run rate, and about 180 basis points across the full year — nearly all of it a one-time charge.

Second, working capital had begun consuming serious cash. In fiscal 2026, receivables grew \$15.4 billion and inventories grew \$11.3 billion, against \$3.1 billion of payables growth and \$7.1 billion of accrued liabilities. Inventories more than doubled to \$21.4 billion, and stood at \$25.8 billion by April 2026. This was either prudent pre-building ahead of the Rubin ramp or the classic late-cycle semiconductor tell, and Raghunathan knew of no way to tell the two apart from outside the company until it was too late to matter.

Third, the balance sheet had quietly become an investment portfolio. As of April 26, 2026, NVIDIA held \$13.2 billion of cash, \$37.1 billion of marketable debt securities, **\$30.2 billion of marketable equity securities**, and **\$43.4 billion of non-marketable securities**, against \$8.5 billion of total debt. Marketable equity holdings more than doubled in a single quarter, from \$12.9 billion at fiscal year-end. In Q1 FY2027, **\$15.9 billion of pre-tax gains on these holdings** flowed through the income statement — which is why GAAP EPS of \$2.39 exceeded non-GAAP EPS of \$1.87. None of it came from selling chips, and Raghunathan struck it out of her starting point before she began.

4.3 Capital Returns

NVIDIA returned \$41.1 billion to shareholders in fiscal 2026 and a record \$20.0 billion in Q1 FY2027 alone. In May 2026 the board added \$80 billion to the buyback authorization and raised the quarterly dividend 25-fold, from \$0.01 to \$0.25 per share.

Shares outstanding are approximately 24.22 billion; weighted-average diluted shares in Q1 FY2027 were 24,391 million.

A 25-fold dividend increase from a company still growing 85% struck Raghunathan as the most interesting disclosure of the quarter, and the least discussed. Read one way, it was confidence in the durability of cash flow. Read another way, it was what companies do when they can no longer reinvest at historical rates of return — and it was usually only obvious which, several years later.

4.4 Market Data as of July 29, 2026

Metric	Value
Share price	\$190.01
Market capitalization	\$4.60 trillion
Shares outstanding	24.22 billion
52-week range	\$164.07 – \$236.54
Trailing twelve-month revenue	\$253.5 billion
Trailing twelve-month EPS (GAAP)	\$6.53
Trailing P/E	29.1×
Forward P/E	19.1×
Beta (5-year, monthly)	2.21
Dividend / yield	\$1.00 / 0.53%
Sell-side consensus target	\$302.83 (61 analysts)
Consensus FY2027 revenue / EPS	\$393.6 billion / \$8.99 (53 analysts)

** The 19.1× is the data provider's next-twelve-months figure, which straddles fiscal 2027 and 2028. On consensus fiscal 2027 EPS of \$8.99 the multiple is 21.1×.*

Either way, the stock traded around 20× forward earnings while growing 85%. On conventional growth-adjusted measures NVIDIA looked *cheap*, and Raghunathan had sat through enough investment committees to know that someone would say so on Monday.

Which was, she thought, exactly the trap. The bear case here had never been that the multiple was too high. It was that the earnings the multiple was applied to might not be durable — and a low forward multiple on a temporary earnings base is not cheapness, it is cheapness-shaped.

5. The Competitive Position

The fund's position had been built on a thesis Raghunathan had written three years earlier and had not revisited in eighteen months: that NVIDIA's moat was software, and that CUDA would hold. She no longer believed the second clause was the important one.

5.1 What NVIDIA Sells

NVIDIA's advantage is not the GPU. It is the vertical integration of seven interlocking product lines shipped as a single rack-scale system: Vera CPU, Rubin GPU, NVLink switch, ConnectX SuperNIC, BlueField DPU, Spectrum Ethernet, and the **Groq LPU** inference accelerator, which NVIDIA obtained through a \$13.0 billion non-exclusive technology license in fiscal 2026. Wrapped around all seven are the CUDA software stack and the Dynamo inference runtime.

A competitor that matches the GPU still has to match the interconnect, the software, and the deployment ecosystem. The 199% growth in networking revenue was, to Raghunathan, the clearest quantitative evidence that customers were buying the system rather than the chip — and the line she would watch first if the moat ever started to leak.

5.2 Who Is Attacking It

Hyperscaler custom silicon. Unit estimates vary widely — one aggregator puts combined 2026 deployment near 1.9 million accelerators (roughly 900,000 Google TPUs, 600,000 AWS Trainium, 250,000 Microsoft Maia, 180,000 Meta MTIA), while some sell-side estimates for Google alone run several times higher. Treat any single number here as an order of magnitude, not a measurement. What is not in dispute: the programs are concentrated in inference, which now represents the majority of AI compute, and they are run by NVIDIA's four largest customers.

AMD. Roughly 5–7% of the data center AI accelerator market by revenue, and rising. Two gigawatt-scale MI450 deployments begin in the second half of 2026 — OpenAI’s, announced in October 2025, and Meta’s, part of a multi-year agreement signed in February 2026. AMD’s accelerators are memory-rich, which matters more in a K3-shaped world than it did in a dense-model world.

Chinese domestic silicon. Constrained by fabrication access but improving, and with a captive market that U.S. export controls have handed to it.

Independent estimates put NVIDIA’s share of data center AI accelerator revenue at **80–85% in 2026, down from approximately 92% in 2023**. The erosion was real, gradual, and — so far — had not touched pricing: gross margin was still 75%. That combination was the thing Raghunathan could not decide how to read. Share loss without price loss is what a widening moat looks like right up until the moment it isn’t, and the four companies taking the share were the same four writing the checks.

5.3 The Demand Side

Four hyperscalers have guided to roughly **\$730 billion of capital expenditure in calendar 2026**, up about 77% from some \$413 billion in 2025: Amazon near \$200 billion, Microsoft near \$190 billion, Alphabet \$195–205 billion (raised on July 22), and Meta \$130–145 billion. Both of the last two were revised upward in the week before this case was written. McKinsey’s base case projects \$5.2 trillion of AI data center capital expenditure cumulatively through 2030, within a scenario range of \$3.7 trillion to \$7.9 trillion.

Against that backdrop, NVIDIA’s roughly \$1 trillion of cumulative Blackwell-plus-Rubin visibility through calendar 2027 was not implausible. It was also not the same thing as a contract, and Raghunathan had made a private rule after 2022 never to discount a number management called “visibility” at the same rate as one it called backlog.

5.4 The Circular Financing Problem

In calendar 2026 alone, NVIDIA has announced more than \$540 billion of investments, guarantees, supply agreements, and partnerships across the AI complex — a figure dominated by a single bidirectional arrangement with SK Group that includes NVIDIA’s own memory purchases, so read it as a measure of entanglement rather than of customer financing. The threads that bear directly on revenue quality:

- A \$30 billion equity stake taken in OpenAI’s early-2026 financing round, superseding a larger non-binding letter of intent from September 2025 that was never converted.
- Up to \$10 billion from NVIDIA and up to \$5 billion from Microsoft committed to Anthropic in November 2025 — caps rather than funded cash — alongside Anthropic’s \$30 billion Azure commitment.
- Under discussion as of late July 2026: guaranteeing up to \$250 billion of OpenAI’s lease payments on a 10-gigawatt SoftBank data center in Ohio, and providing up to \$350 billion of financing for OpenAI GPU purchases. These talks are early and may not close. For scale, NVIDIA’s Q1 FY2027 10-Q discloses maximum gross exposure under *all* existing facility lease guarantees of \$3.5 billion — the proposal is roughly seventy times the entire disclosed book.

This was the part that genuinely worried her, more than K3 did. When a supplier funds its customers’ purchases, reported revenue growth is partly a financing decision rather than an independent measure of demand. Ordinarily a customer’s willingness to pay is evidence that the product creates value; when the supplier provides the money, that evidence weakens. It is also the mechanism by which a hardware vendor’s balance sheet absorbs its customers’ credit risk, and several of the largest AI infrastructure borrowers — OpenAI, SoftBank, CoreWeave, Nebius, Oracle — had materially increased leverage since 2025.

None of that made the revenue fake; OpenAI genuinely used the machines. But it meant some fraction of the forecast depended on counterparties NVIDIA was itself financing, and Raghunathan could see no principled reason those cash flows deserved the same discount rate as cash flows from Microsoft. A credit event at any one of them would hit the position through three channels at once — lost revenue, impaired equity stakes, called guarantees — and the three were perfectly correlated.

6. The Valuation Question

6.1 Framing It Backwards

Raghunathan’s first instinct was always to forecast forward, and she had trained herself out of it. A forward forecast is only as good as its weakest assumption, and it is far too easy to bury a heroic one inside a growth rate where nobody — including its author — will notice it again. Working backward from the market price forces the heroic assumption into the open. It was

the approach Aswath Damodaran had long advocated for large-cap technology valuations, and it had the additional virtue, in an investment committee, of changing the question from *what do you think it is worth* to *what would have to be true*. The logic:

1. Start with the market’s enterprise value.
2. Impose realistic steady-state assumptions about margins, reinvestment, taxes, and cost of capital.
3. Solve for the revenue required to justify that enterprise value.
4. Ask three questions of that revenue number: **Is it possible? Is it plausible? Is it probable?**

Step one. At \$190.01 and 24.22 billion shares, equity value is \$4.602 trillion. Non-operating assets — cash, marketable debt and equity securities, and non-marketable securities, net of \$8.5 billion of debt — total approximately \$115 billion. Enterprise value is therefore approximately **\$4.49 trillion**.

The \$43.4 billion of non-marketable securities within that total consists largely of stakes in private AI companies whose value is correlated with the very demand NVIDIA is selling into.

Steps two and three. Suppose NVIDIA were already a mature company converting revenue to free cash flow at a steady rate, growing at 3% in perpetuity. The required annual free cash flow is enterprise value times (WACC – g):

WACC	Required FCF	Revenue at 35% FCF margin	at 42%	at 50%
9.5%	\$292B	\$833B	\$694B	\$583B
11.0%	\$359B	\$1,026B	\$855B	\$718B
12.5%	\$426B	\$1,218B	\$1,015B	\$852B

Step four. For context: NVIDIA’s trailing twelve-month revenue is \$253 billion. The **entire global semiconductor industry** was approximately \$795 billion in calendar 2025 and is forecast at \$1.51 trillion in 2026, with \$1.9 trillion projected for 2027 — a forecast that itself embeds an unprecedented memory supercycle.

So the answer to “is it possible?” was yes — but the number it took was uncomfortable. At a 42% steady-state free cash flow margin and an 11% cost of capital, today’s price required NVIDIA — one company, in one segment of one industry — to reach a steady-state revenue run rate of roughly **\$855 billion**, more than the entire semiconductor industry generated in 2025.

Raghuathan wrote the figure down and sat with it. It was not absurd. Hyperscaler capital expenditure alone was guided to \$730 billion for the year, and the industry had beaten nearly every forecast made about it since 2023. But it was the kind of number that, stated plainly, tends to lose arguments it should win — and she had noticed that the people quoting NVIDIA’s 19× forward multiple in her Monday meetings were never the same people who had run this calculation.

6.2 The Big Market Delusion

Bradford Cornell and Aswath Damodaran’s phrase for the systematic error here is the **big market delusion**: when a market is genuinely enormous, every participant assumes it will capture a large share of it, and the sum of the implied shares far exceeds 100%.

The arithmetic is straightforward. Take the market capitalizations of NVIDIA (\$4.60T), Broadcom (\$1.76T), TSMC (\$1.76T), and AMD (\$0.70T); add the AI-attributable portion of the hyperscalers’ valuations and the private marks on OpenAI, Anthropic, and the neoclouds. The aggregate AI revenue required to justify that total, at any plausible aggregate margin, can then be compared to credible forecasts of AI end-market spending.

The test does not identify NVIDIA specifically as overvalued. It establishes that the sector’s *aggregate* implied revenue is likely inconsistent, and therefore that at least some of the companies priced for dominance will not achieve it. NVIDIA, Raghuathan thought, had the strongest claim of any of them to being the winner. The uncomfortable question was whether that claim was already fully in the price.

6.3 The Kimi K3 Question, Stated Precisely

By Thursday evening Raghuathan had reduced the whole controversy to one question with three defensible answers. She could construct a coherent world under each. What she could not do was hold all three at once and still produce a number.

Answer A — Deflationary (bearish). Efficiency gains are real and compounding. K3 delivers roughly 2.5× better scaling efficiency than its predecessor and near-frontier capability at a fraction of the price. If capability per dollar of hardware keeps improving at this rate, the compute required to deliver any fixed level of AI service falls fast. Hyperscalers eventually recognize

they have over-provisioned; the \$730 billion capex cohort cuts; NVIDIA's revenue does not just decelerate, it declines. Open weights accelerate this by collapsing the pricing power of the frontier labs that are NVIDIA's most aggressive buyers.

Answer B — Jevons (bullish). This is Huang's argument and the sell-side consensus. When a resource becomes cheaper to use, total consumption rises rather than falls. Cheap, open, near-frontier intelligence unlocks applications that were previously uneconomic — agentic workflows, always-on inference, every enterprise self-hosting a model on its own data. Aggregate token demand rises faster than efficiency improves. UBS drew the analogy to mobile telecommunications: each generation delivered dramatically cheaper data per bit and dramatically higher total network investment.

Answer C — Mix shift (bullish on volume, bearish on margin). Demand rises, but its composition changes. K3-shaped workloads are memory-bound and inference-heavy. Inference is more price-sensitive than training, more amenable to custom silicon, and less dependent on CUDA. NVIDIA sells more systems but at lower margin, and the value migrates toward memory suppliers and packaging capacity.

7. Friday

The three answers sat on Raghunathan's whiteboard in her own handwriting, and by Thursday night she had argued herself into each of them at least once.

Every piece of evidence she trusted pointed in two directions at the same time. Revenue was accelerating — 56%, then 62%, then 73%, then 85% year over year — which was not what a business absorbing a demand shock looks like, and not what a maturing one looks like either. But the acceleration had been financed in part by NVIDIA itself, and she did not know how much. The memory market was forecast to grow 250% in a year, which confirmed that inference demand was real and simultaneously suggested that the value was migrating toward suppliers NVIDIA had to buy from. Gross margin was still 75%, which meant the share NVIDIA had already lost had cost it nothing — or meant that the price war had not started yet.

She had a week's worth of facts and no way to make them contradict each other less.

What she did know was the shape of the problem. A world in which growth decelerates hard and margins converge toward semiconductor norms would value the shares at a fraction of \$190. A world in which inference compounds for another decade and the rack-scale moat holds would value them at a large multiple of it. The distance between those two worlds was wider than for anything else she covered at this size, and it was not obvious to her that the market had picked the wrong point between them — only that it had picked one, and moved more than 10% in ten sessions without a single new financial disclosure.

The memo was due Friday. The committee met Monday. The CIO had not asked for conviction: he had asked what the business was worth, whether the market was paying more or less than that, and what would have to be true for each answer to be right.

Raghunathan opened a blank model and started with the only number she was sure of: fiscal 2026 revenue of \$215.9 billion.

8. Discussion Questions

Valuation methodology

1. In most discounted cash flow models of high-growth companies, terminal value accounts for 60–80% of total value. What does that imply about how much of any NVIDIA valuation is really a forecast of the 2040s, and how should it affect the confidence anyone places in a point estimate?
2. NVIDIA holds roughly \$115 billion in non-operating assets, of which \$43 billion is non-marketable stakes in private AI companies it also sells to. How should these be treated when bridging from enterprise value to equity value? Does adding them at carrying value double-count anything?
3. NVIDIA's five-year beta is 2.21, implying a cost of equity above 14% at current risk-free rates and equity risk premiums. A bottom-up beta — one built from comparable semiconductor companies rather than regressed from NVIDIA's own returns — would be closer to 1.35–1.75. Which is right for a company with 75% gross margins, no net debt, and extreme demand cyclicity? How much does the choice matter?

The Kimi K3 question

4. Raghunathan could not choose among Answers A, B, and C in Section 6.3. Which dominates over a five-year horizon? Identify the evidence in the case that most strongly supports your choice, and the evidence that most strongly cuts against it — then say which assumptions, growth or margin or both, would carry that view into a valuation.
5. In January 2025, DeepSeek’s release cost NVIDIA \$589 billion of market value in one day, and the stock fully recovered within six months. Does that history make the July 2026 selloff a buying opportunity, or is reasoning by analogy from a single precedent the exact error that destroys portfolios? What is different this time, if anything?
6. Kimi K3 is memory-defined rather than compute-defined. If the industry’s models converge on that shape, is NVIDIA advantaged or disadvantaged relative to AMD, Broadcom, and the memory manufacturers? Raghunathan could argue it both ways — resolve her indecision.

Competition and structure

7. NVIDIA’s four largest customers are also its four most credible competitors. What path of NVIDIA market share is defensible through 2030, and why might share lost to a customer’s in-house silicon be more damaging than share lost to AMD?
8. Raghunathan concluded that revenue from counterparties NVIDIA is financing does not deserve the same discount rate as revenue from Microsoft. Is she right? Under what circumstances does vendor financing make a supplier’s revenue *lower* quality than a peer’s, and how would an outside analyst measure the exposure?
9. Apply the big market delusion test of Section 6.2. Summing the AI-attributable market capitalization across the sector and deriving the aggregate revenue it requires, is the sector internally consistent? If not, which companies are mispriced, and where does NVIDIA sit among them?

Judgment

10. What should Raghunathan recommend on Monday — buy, hold, or reduce — and at what price would the answer flip? Which single observable metric over the next four quarters would most change her view, at what threshold, and in which direction?

9. Sources

Company filings and disclosures

- NVIDIA Corporation, “NVIDIA Announces Financial Results for Fourth Quarter and Fiscal 2026,” February 25, 2026. <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-fourth-quarter-and-fiscal-2026>
- NVIDIA Corporation, “NVIDIA Announces Financial Results for First Quarter Fiscal 2027,” May 20, 2026. <https://nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-first-quarter-fiscal-2027>
- NVIDIA Corporation, Form 8-K (Q4 and fiscal 2026 results), February 25, 2026. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000008/nvidia-20260225.htm>
- NVIDIA Corporation, Form 8-K (Q1 fiscal 2027 results), May 20, 2026. <https://www.sec.gov/Archives/edgar/data/1045810/000104581026000010/nvidia-20260520.htm>
- NVIDIA Corporation, Investor Relations. <https://investor.nvidia.com/>
- NVIDIA Corporation, “NVIDIA Vera Rubin Opens Agentic AI Frontier,” March 16, 2026. <https://nvidianews.nvidia.com/news/nvidia-vera-rubin-platform>

Kimi K3

- Moonshot AI, “Kimi K3: Open Frontier Intelligence” (technical blog). <https://www.kimi.com/blog/kimi-k3>
- Moonshot AI, Kimi K3 model card and Kimi K3 License. <https://huggingface.co/moonshotai/Kimi-K3>
- Fortune, “Markets experience new DeepSeek shock after MoonShot AI releases Kimi K3,” July 17, 2026. <https://fortune.com/2026/07/17/china-moonshot-kimi-k3-markets-china-ai/>
- Tom’s Hardware, “China’s 2.8-trillion-parameter Kimi K3 beats Claude Fable 5 in Frontend Code Arena benchmark,” July 2026. <https://www.tomshardware.com/tech-industry/artificial-intelligence/moonshot-releases-2-8-trillion-parameter-kimi-k3>
- Nathan Lambert, “Kimi K3: The open-weights escalation,” *Interconnects*, July 2026. <https://www.interconnects.ai/p/kimi-k3-the-open-weights-escalation>
- CNN Business, “What is China’s Kimi K3 and why is the US so rattled by it?,” July 23, 2026. <https://www.cnn.com/2026/07/23/tech/china-ai-moonshot-kimi-explainer-intl-hnk>

- Tom's Hardware, "China's Moonshot AI reportedly used Nvidia Blackwell chips for training Kimi K3," July 2026. <https://www.tomshardware.com/tech-industry/artificial-intelligence/chinas-moonshot-ai-reportedly-used-nvidia-blackwell-chips-for-training-kimi-k3>
- Axios, "Exclusive: Nvidia's Jensen Huang defends Chinese AI amid Kimi panic," July 22, 2026. <https://www.axios.com/2026/07/22/nv-jensen-huang-china-open-source-ai>

Market reaction and analyst commentary

- Bloomberg, "US Stocks Rebound From Selloff as Nvidia, Other Chipmakers Rise," July 20, 2026. <https://www.bloomberg.com/news/articles/2026-07-20/us-stocks-rebound-from-selloff-as-nvidia-other-chipmakers-rise>
- Bloomberg, "Moonshot's Kimi K3 May Be More About Memory Than Compute," July 20, 2026. <https://www.bloomberg.com/news/articles/2026-07-20/moonshot-s-kimi-k3-may-be-more-about-memory-than-compute>
- Axios, "Nvidia reignites 'circular' AI concerns as it weighs OpenAI financing guarantee," July 27, 2026. <https://www.axios.com/2026/07-27/nvidia-reignites-circular-ai-concerns-as-it-weighs-openai-financing-guarantee>
- Benzinga, "Nvidia Funds AI Frenzy: Timeline of its Circular Financing Deals," July 2026. <https://www.benzinga.com/markets/tech/26/07/nvidia-funds-ai-frenzy-timeline-of-its-circular-financing-deals-so-far>
- Stock Analysis, "NVIDIA (NVDA) Stock Price & Overview." <https://stockanalysis.com/stocks/nvda/>

Industry data

- World Semiconductor Trade Statistics, Spring 2026 forecast. <https://www.wsts.org/76/Recent-News-Release>
- McKinsey & Company, "The cost of compute: A \$7 trillion race to scale data centers." <https://www.mckinsey.com/industries/technology-media-and-telecommunications/our-insights/the-cost-of-compute-a-7-trillion-dollar-race-to-scale-data-centers>
- Data Center Knowledge, "GTC 2026: Nvidia Unveils Vera Rubin AI Platform, Eyes \$1T by 2027." <https://www.datacenterknowledge.com/center-chips/gtc-2026-nvidia-unveils-vera-rubin-ai-platform-eyes-1t-by-2027>
- AMD, Advancing AI 2026 keynote, July 23, 2026 (source of the \$1.4 trillion 2030 AI accelerator TAM estimate).
- Alphabet Q2 2026 earnings, July 22, 2026 (2026 capital expenditure guidance raised to \$195–205 billion).

Valuation framework

- Aswath Damodaran, *Musings on Markets*. <https://aswathdamodaran.blogspot.com/>
- Aswath Damodaran, "Equity Risk Premiums (ERP): Determinants, Estimates and Implications — The 2026 Edition." https://papers.ssrn.com/sol3/papers.cfm?abstract_id=6361419
- Aswath Damodaran, *Narrative and Numbers: The Value of Stories in Business*. Columbia Business School Publishing.
- Bradford Cornell and Aswath Damodaran, "The Big Market Delusion: Valuation and Investment Implications," *Financial Analysts Journal* (2020).

Exhibits

Fiscal calendar. NVIDIA's fiscal year ends in late January. Fiscal 2026 ran from January 27, 2025 to January 25, 2026. Fiscal 2027 began January 26, 2026 and ends in late January 2027. All per-share figures are stated on a post-split basis (10-for-1 split, June 2024).

Financial data source: NVIDIA Forms 8-K of February 25, 2026 (Q4 and fiscal 2026) and May 20, 2026 (Q1 fiscal 2027), and the related SEC filings. **Market data:** closing prices through July 29, 2026.

Exhibit 1 — Consolidated Statements of Income

(US\$ millions, except per-share data)

Fiscal year ended	FY2023	FY2024	FY2025	FY2026
Revenue	26,974	60,922	130,497	215,938
Cost of revenue	11,618	16,621	32,639	62,475
Gross profit	15,356	44,301	97,858	153,463
<i>Gross margin</i>	<i>56.9%</i>	<i>72.7%</i>	<i>75.0%</i>	<i>71.1%</i>
Total operating expenses ¹	11,132	11,329	16,405	23,076
<i>Operating expenses / revenue</i>	<i>41.3%</i>	<i>18.6%</i>	<i>12.6%</i>	<i>10.7%</i>
Operating income	4,224	32,972	81,453	130,387
<i>Operating margin</i>	<i>15.7%</i>	<i>54.1%</i>	<i>62.4%</i>	<i>60.4%</i>
Interest income	267	866	1,786	2,300
Interest expense	(262)	(257)	(247)	(259)
Other income (expense), net	(48)	237	1,034	9,022
Income before income taxes	4,181	33,818	84,026	141,450
Income tax expense (benefit)	(187)	4,058	11,146	21,383
<i>Effective tax rate</i>	<i>n.m.</i>	<i>12.0%</i>	<i>13.3%</i>	<i>15.1%</i>
Net income	4,368	29,760	72,880	120,067
<i>Net margin</i>	<i>16.2%</i>	<i>48.8%</i>	<i>55.8%</i>	<i>55.6%</i>
Diluted EPS	\$0.17	\$1.19	\$2.94	\$4.90
Non-GAAP diluted EPS	—	—	\$2.99	\$4.77

¹ Research and development plus sales, general and administrative. FY2023 also includes a \$1,353 million acquisition-termination charge (the abandoned Arm transaction). Detailed R&D/SG&A splits are in the 10-K.

Exhibit 2 — Quarterly Results, FY2026 through Q1 FY2027

2A. Reported results

(US\$ millions)

Line item	Q1 FY26	Q2 FY26	Q3 FY26	Q4 FY26	FY2026	Q1 FY27
Quarter ended	Apr 27 '25	Jul 27 '25	Oct 26 '25	Jan 25 '26		Apr 26 '26
Revenue	44,062	46,743	57,006	68,127	215,938	81,615
Sequential growth	—	+6.1%	+22.0%	+19.5%		+19.8%
Year-over-year growth	+69%	+56%	+62%	+73%	+65%	+85%
Gross profit	26,668	33,853	41,849	51,093	153,463	61,157
<i>GAAP gross margin</i>	<i>60.5%</i>	<i>72.4%</i>	<i>73.4%</i>	<i>75.0%</i>	<i>71.1%</i>	<i>74.9%</i>
Operating income	21,638	28,440	36,010	44,299	130,387	53,536

Line item	Q1 FY26	Q2 FY26	Q3 FY26	Q4 FY26	FY2026	Q1 FY27
<i>Operating margin</i>	<i>49.1%</i>	<i>60.8%</i>	<i>63.2%</i>	<i>65.0%</i>	<i>60.4%</i>	<i>65.6%</i>
Net income (GAAP)	18,775	26,422	31,910	42,960	120,067	58,321
Diluted EPS (GAAP)	\$0.76	\$1.08	\$1.30	\$1.76	\$4.90	\$2.39
Diluted EPS (non-GAAP)	\$0.78	—	—	\$1.62	\$4.77	\$1.87

Q1 FY2026 cost of revenue includes a \$4,500 million charge on H2O inventory affected by U.S. export controls. Q4 FY2026 Data Center revenue was \$62.3 billion. In Q1 FY2027, other income included \$15,936 million of pre-tax gains on equity securities, lifting income before tax to \$69,903 million.

2B. Company outlook for Q2 fiscal 2027, issued May 20, 2026

Item	Guidance
Revenue	\$91.0 billion, $\pm$ 2%
GAAP gross margin	74.9%, $\pm$ 50 bps
Non-GAAP gross margin	75.0%, $\pm$ 50 bps
GAAP operating expenses	~\$8.5 billion
Non-GAAP operating expenses	~\$8.3 billion
Effective tax rate, full year FY2027	16–18%
China Data Center compute revenue	none assumed

Exhibit 3 — Revenue by Segment

(US\$ millions)

3A. Annual, under the reporting structure through FY2026

Fiscal year	Data Center	Gaming	Professional Visualization	Automotive	OEM & Other	Total
FY2023	15,005	9,067	1,544	903	455	26,974
<i>% of total</i>	<i>55.6%</i>	<i>33.6%</i>	<i>5.7%</i>	<i>3.3%</i>	<i>1.7%</i>	
FY2024	47,525	10,447	1,553	1,091	306	60,922
<i>% of total</i>	<i>78.0%</i>	<i>17.1%</i>	<i>2.5%</i>	<i>1.8%</i>	<i>0.5%</i>	
<i>YoY growth</i>	<i>+217%</i>	<i>+15%</i>	<i>+1%</i>	<i>+21%</i>	<i>−33%</i>	<i>+126%</i>
FY2025	115,186	11,350	1,878	1,694	389	130,497
<i>% of total</i>	<i>88.3%</i>	<i>8.7%</i>	<i>1.4%</i>	<i>1.3%</i>	<i>0.3%</i>	
<i>YoY growth</i>	<i>+142%</i>	<i>+9%</i>	<i>+21%</i>	<i>+55%</i>	<i>+27%</i>	<i>+114%</i>
FY2026	193,737	16,042	3,191	2,349	619	215,938
<i>% of total</i>	<i>89.7%</i>	<i>7.4%</i>	<i>1.5%</i>	<i>1.1%</i>	<i>0.3%</i>	
<i>YoY growth</i>	<i>+68%</i>	<i>+41%</i>	<i>+70%</i>	<i>+39%</i>	<i>+59%</i>	<i>+65%</i>

3B. Q1 FY2027, under the current reporting structure

Beginning in fiscal 2027 NVIDIA reports two market platforms and discloses a compute/networking split within Data Center.

Market platform	Q1 FY2027	YoY growth	Sequential growth
Data Center	75,246	+92%	+21%
of which: Compute	~60,400	+77%	+18%
of which: Networking	~14,800	+199%	+35%
Edge Computing²	6,369	+29%	+10%
Total revenue	81,615	+85%	+20%

Compute and networking are the rounded figures NVIDIA disclosed under its previous sub-market definitions; they do not sum exactly to the Data Center total.

² In NVIDIA's words, Edge Computing "highlights data processing devices for agentic and physical AI including PCs, game consoles, workstations, AI-RAN base stations, robotics and automotive" — consolidating the former Gaming, Professional Visualization, and Automotive lines.

For fiscal 2026 as a whole, Data Center split into **\$162,361 million of compute** and **\$31,376 million of networking**.

Data Center demand mix, Q1 FY2027

Customer type	Revenue	Share of Data Center
Hyperscale	37,869	50.3%
AI Clouds, Industrial and Enterprise ("ACIE")	37,377	49.7% (up 31% sequentially)

Exhibit 4 — Consolidated Balance Sheets

(US\$ millions)

Line item	Jan 26, 2025 (FY25)	Jan 25, 2026 (FY26)	Apr 26, 2026 (Q1 FY27)
Assets			
Cash and cash equivalents	8,589	10,605	13,237
Marketable debt securities	34,621	39,065	37,098
Marketable equity securities	—	12,886	30,237
<i>Cash and marketable securities (total)</i>	<i>43,210</i>	<i>62,556</i>	<i>80,572</i>
Accounts receivable, net	23,065	38,466	40,710
Inventories	10,080	21,403	25,797
Prepaid expenses and other	3,771	3,180	3,916
Total current assets	80,126	125,605	150,995
Property and equipment, net	6,283	10,383	12,403
Operating lease assets	1,793	2,867	4,258
Goodwill	5,188	20,832	20,894
Intangible assets, net	807	3,306	3,120
Deferred income tax assets	10,979	13,258	11,707
Non-marketable securities	3,387	22,251	43,364
Other assets	3,038	8,301	12,733
Total assets	111,601	206,803	259,474
Liabilities			
Accounts payable	6,310	9,812	13,097
Accrued and other current liabilities	11,737	21,352	29,787
Short-term debt	—	999	1,000
Total current liabilities	18,047	32,163	43,884
Long-term debt	8,463	7,469	7,470
Long-term operating lease liabilities	1,519	2,572	3,878
Other long-term liabilities	4,245	7,306	8,768
Total liabilities	32,274	49,510	64,000
Total shareholders' equity	79,327	157,293	195,474

Derived metrics

Metric	FY2025	FY2026	Q1 FY2027
Total debt	8,463	8,468	8,470
Net cash (cash + marketable debt securities – debt)	34,747	41,202	41,865
All non-operating assets net of debt³	38,134	76,339	115,466

Metric	FY2025	FY2026	Q1 FY2027
Debt / equity	10.7%	5.4%	4.3%
Days sales outstanding [□]	65	65	45
Days inventory outstanding [□]	113	125	115
Return on equity [□]	91.9%	76.3%	119.3%

³ Cash + marketable debt securities + marketable equity securities + non-marketable securities – total debt. At Q1 FY2027 this equals \$4.77 per share.

[□] Annual columns use a 365-day year; the Q1 FY2027 column uses the 91-day quarter. The two bases are not directly comparable.

[□] Net income divided by *ending* shareholders' equity; Q1 FY2027 is annualized. Q1 FY2027 net income includes \$15.9 billion of pre-tax gains on equity securities.

Exhibit 5 — Cash Flow and Capital Allocation

(US\$ millions)

Line item	FY2024	FY2025	FY2026	Q1 FY2027
Net income	29,760	72,880	120,067	58,321
Stock-based compensation	3,549	4,737	6,386	1,928
Depreciation and amortization	—	1,864	2,843	997
Gains on equity securities, net	—	(1,030)	(8,918)	(15,936)
Deferred income taxes	—	(4,477)	(1,424)	1,584
Change in accounts receivable	—	(13,063)	(15,399)	(2,243)
Change in inventories	—	(4,781)	(11,324)	(4,420)
Change in accounts payable	—	3,357	3,096	2,210
Change in accrued and other liabilities	—	5,499	7,101	7,763
Cash from operating activities	28,090	64,089	102,718	50,344
Purchases of property, equipment and intangibles	(1,069)	(3,236)	(6,042)	(1,757)
Principal payments on property, equipment and intangibles	n/a	(129)	(101)	(33)
Free cash flow[□]	27,021	60,724	96,575	48,554
<i>FCF margin</i>	<i>44.4%</i>	<i>46.5%</i>	<i>44.7%</i>	<i>59.5%</i>
Purchases of non-marketable securities	—	(1,486)	(17,502)	—
Groq licensing payment	—	—	(13,000)	—
Share repurchases	(9,533)	(33,706)	(40,086)	(19,312)
Dividends paid	(395)	(834)	(974)	—
Total shareholder returns	9,928	34,540	41,060	~20,000
<i>Payout as % of FCF</i>	<i>37%</i>	<i>57%</i>	<i>43%</i>	<i>41%</i>

[□] NVIDIA defines free cash flow as operating cash flow less *both* purchases of and principal payments on property, equipment and intangible assets. The FY2024 figure is operating cash flow less capital expenditure only. The \$13.0 billion Groq payment was consideration for a non-exclusive license to Groq's LPU inference technology, not an equity investment.

Capital intensity

Line item	FY2024	FY2025	FY2026	Q1 FY2027
Capex / revenue	1.8%	2.5%	2.8%	2.2%
D&A / revenue	—	1.4%	1.3%	—
Working capital absorption [□]	—	12.9%	19.3%	n.m.

[□] Net change in operating working capital (receivables plus inventories, less payables and accrued liabilities) as a percentage of the year's revenue *increase*. In Q1 FY2027 the change was negative — working capital released cash, as a \$7.8 billion build in accrued liabilities exceeded higher receivables and inventory — so no quarterly figure is shown.

Capital returns policy, as of May 2026

- Quarterly dividend raised from \$0.01 to **\$0.25 per share** — a 25-fold increase.
- **\$80 billion** added to the share repurchase authorization, on top of approximately \$39 billion remaining.
- A record \$20.0 billion returned in Q1 FY2027 alone.

Exhibit 6 — Market Data and Valuation Multiples

6A. NVIDIA, as of the close on July 29, 2026

Metric	Value
Share price	\$190.01
Shares outstanding	24.22 billion
Weighted-average diluted shares (Q1 FY27)	24,391 million
Market capitalization	\$4.602 trillion
Less: non-operating assets net of debt	\$115.5 billion
Implied enterprise value	\$4.487 trillion
52-week range	\$164.07 – \$236.54
Trailing twelve-month revenue	\$253.5 billion
Trailing twelve-month GAAP EPS	\$6.53
Beta (5-year monthly regression)	2.21
Beta — bottom-up semiconductor comparables	1.35 – 1.75
10-year U.S. Treasury yield, July 30, 2026	4.70%
Implied equity risk premium (Damodaran, January 2026)	4.23%
Annual dividend / yield	\$1.00 / 0.53%
Sell-side consensus price target	\$302.83 (61 analysts)
Consensus FY2027 revenue	\$393.6 billion (53 analysts)
Consensus FY2027 EPS	\$8.99

6B. Implied multiples

Multiple	On FY2026 actuals	On TTM	On consensus FY2027E
EV / Revenue	20.8×	17.7×	11.4×
EV / Operating income	34.4×	—	~17–18×
P / E	38.8×	29.1×	21.1×
P / Free cash flow	47.7×	—	—

6C. Peer comparison, July 2026

Company	Market cap	Trailing P/E	Forward P/E	Position in AI stack
NVIDIA	\$4.60T	29.1×	19.1×	Accelerated computing systems
TSMC	\$1.76T	25.3×	17.8×	Foundry; manufactures for all
Broadcom	\$1.76T	61.6×	23.5×	Custom ASICs and networking silicon

Company	Market cap	Trailing P/E	Forward P/E	Position in AI stack
AMD	\$0.70T	143.2×	48.2×	Merchant GPU alternative

All four as of the July 29, 2026 close, from the same data provider: Forward multiples are next-twelve-months.

The FY2027E column uses consensus revenue of \$393.6 billion and EPS of \$8.99. Data providers quote a $19.1\times$ “forward P/E” on a next-twelve-months basis that straddles fiscal 2027 and 2028; on fiscal 2027 consensus the multiple is $21.1\times$.

Exhibit 7 — Price Path Through the July 2026 Selloff

7A. Where the stock stands, July 29, 2026

Metric	Value
Close, July 29, 2026	\$190.01
Market capitalization	\$4.60 trillion
Shares outstanding	24.22 billion
52-week range	\$164.07 – \$236.54
Distance from 52-week high	−19.7%
Distance from 52-week low	+15.8%
Decline since July 15, 2026	−10.6% (≈ \$545 billion of market value)

7B. Daily closes through the selloff

Date (2026)	Close	Change	Context
Wed Jul 15	\$212.50		Pre-event reference
Thu Jul 16	\$207.40	−2.4%	Moonshot AI unveils Kimi K3
Fri Jul 17	\$202.81	−2.2%	Global semiconductor selloff; K3 exhibited on the opening day of the World AI Conference in Shanghai. The SOX ends the week down roughly 10% — its worst week since April 2025 — and enters a bear market
Mon Jul 20	\$203.28	+0.2%	Broad rebound; SOX +0.6% as banks publish “buy the dip” notes
Tue Jul 21	\$207.29	+2.0%	Recovery continues
Wed Jul 22	\$212.06	+2.3%	Full round trip. White House accuses Moonshot of training K3 on NVIDIA GB300 systems in Thailand
Thu Jul 23	\$208.76	−1.6%	
Fri Jul 24	\$206.84	−0.9%	
Mon Jul 27	\$196.51	−5.0%	Moonshot publishes full K3 weights on Hugging Face. Reports that NVIDIA may guarantee up to \$250B of OpenAI lease payments and finance up to \$350B of GPU purchases; circular-financing concerns resurface
Tue Jul 28	\$197.01	+0.3%	
Wed Jul 29	\$190.01	−3.6%	Case valuation date

Reference point: June 30, 2026 close, \$200.09.

7C. Sector context

Index / measure	Move
Philadelphia Semiconductor Index (SOX), week ending July 17	–10.0%, worst week since April 2025; index closed more than 20% below its late-June record
SOX, July 20	+0.6%
SOX, prior twelve months through early July 2026	approximately +130%
NVIDIA market value lost, week ending July 17	approximately \$197 billion

7D. Published analyst response

Notes published in the week following the release:

- **Bank of America** (Vivek Arya’s team) argued the low price of Chinese open models reflects a business-model choice rather than a collapse in hardware cost, and that enterprise deployment of open models *increases* demand for HBM, DRAM, and NAND. The team noted Chinese open-source APIs price at 5× to 350× below Western equivalents.
- **UBS** (Sunil Tirumalai; Ulrike Hoffmann-Burchardi) estimated Chinese models handle routine tasks at 15–20% of the cost of comparable U.S. models, but argued hyperscalers treat AI as winner-takes-all and will not cut capital expenditure. UBS drew an analogy to the 2G-to-5G transition: each generation delivered radically cheaper data and radically higher total network investment.
- **Morgan Stanley** called the memory selloff an entry opportunity given persistent data center memory shortages.
- **Jensen Huang** called the Chinese models “excellent,” said that “open-source models that are excellent should be used,” and summarized his read of the selloff as “free AI should be great for hardware.” On July 23 he published an open letter, *Open Weights and American AI Leadership*, opposing proposals to restrict downloadable AI models in the United States.

Exhibit 8 — Competitive and Industry Context

8A. Data center AI accelerator market share (revenue basis, estimated)

Vendor	2023	2026E
NVIDIA	~92%	75–85%
AMD	~2%	5–7%
Hyperscaler custom silicon	~4%	8–12%
Other (Intel, startups, Chinese domestic)	~2%	2–4%

Independent industry-tracker estimates. Ranges do not sum to 100%; methodologies differ on whether custom ASICs belong in the denominator. NVIDIA’s 2026 share is quoted from 70% to 86% depending on that choice. The 2023 figure is a data center GPU number; some trackers put it at 98% on a revenue basis.

8B. Hyperscaler custom silicon deployment, 2026 (estimated units)

Program	Owner	Est. 2026 accelerators	Primary workload
TPU v7 “Ironwood”	Google	~900,000	Internal + Google Cloud
Trainium 3	Amazon	~600,000	Internal + AWS inference
Maia 200	Microsoft	~250,000	Internal Azure
MTIA	Meta	~180,000	Internal ranking + inference
Total		~1,930,000	

Single-source aggregator estimates. Definitions differ across trackers: the Google figure combines two TPU generations and the Amazon figure combines Trainium and Inferentia. Some sell-side estimates for Google alone are several times higher.

These programs are run by NVIDIA’s four largest customers, are concentrated in inference (now the majority of AI compute), and are growing faster than the overall market. Separately, two gigawatt-scale AMD MI450 deployments begin in the second half of 2026: OpenAI’s, announced October 2025, and Meta’s, under a multi-year agreement signed February 2026.

8C. Announced 2026 capital expenditure, four largest hyperscalers

Company	2026 capex guidance
Amazon	~\$200 billion
Microsoft	~\$190 billion
Alphabet	\$195–205 billion (raised July 22, 2026 from \$180–190 billion)
Meta	\$130–145 billion (low end raised July 2026)
Combined	~\$730 billion (up ~77% from ~\$413 billion in 2025)

8D. Market size context

Measure	Figure
Global semiconductor market, calendar 2025	~\$795 billion
Global semiconductor market, 2026 forecast (WSTS Spring 2026)	\$1.51 trillion (+90%)
Global semiconductor market, 2027 forecast	~\$1.9 trillion (+27% on WSTS's unrounded figures)
of which memory, 2026	>\$800 billion (+~250%)
Cumulative AI data center capex through 2030 (McKinsey base case)	\$5.2 trillion (range \$3.7–7.9 trillion)
AI accelerator market in 2030 (AMD, July 23, 2026)	\$1.4 trillion , raised from \$500 billion a year earlier
NVIDIA cumulative Blackwell + Rubin revenue visibility through CY2027	~\$1 trillion (raised from \$500 billion)
NVIDIA trailing twelve-month revenue	\$253.5 billion

8E. Technology roadmap

Vendor	Current generation	Next
NVIDIA	Blackwell (B200/B300/GB300); Vera Rubin announced in full production on March 16, 2026 , with partner systems available in the second half of the year	Rubin Ultra, then Feynman
AMD	MI355X (288 GB HBM3E)	MI450; gigawatt-scale OpenAI and Meta deployments begin H2 2026
Google	TPU v7 “Ironwood”	Annual cadence
Amazon	Trainium 3	Trainium 4
Microsoft	Maia 200	—

Vera Rubin platform specifications

Specification	VR200 accelerator	NVL72 rack
Compute	50 PFLOPS (FP4)	3.6 exaflops
Memory	288 GB	—
Memory bandwidth	22 TB/s	260 TB/s all-to-all NVLink

The platform ships as seven co-designed products: Vera CPU, Rubin GPU, NVLink 6 switch, ConnectX-9 SuperNIC, BlueField-4 DPU, Spectrum-6 Ethernet, and the Groq 3 LPU inference accelerator. (NVIDIA described six chips in February 2026; the Groq LPU was folded into the platform at GTC in March.)

8F. Counterparty and financing exposure

Arrangement	Scale	Status as of July 30, 2026
OpenAI equity stake	\$30 billion	Closed in the company's early-2026 round, superseding a larger non-binding LOI from September 2025 that was never converted
Anthropic investment	up to \$10 billion (NVIDIA) plus up to \$5 billion (Microsoft)	Announced November 2025; caps, not funded cash
Groq — non-exclusive technology license	\$13.0 billion	Closed in FY2026; Groq LPU now shipping inside Vera Rubin. Not an equity investment

Arrangement	Scale	Status as of July 30, 2026
OpenAI data center lease guarantee (SoftBank, Ohio, 10 GW)	up to \$250 billion	Under discussion; may not close. The Q1 FY2027 10-Q discloses \$3.5 billion of maximum gross exposure under all existing facility lease guarantees
OpenAI GPU purchase financing	up to \$350 billion	Under discussion; may not close
Total announced 2026 commitments across the AI complex	>\$540 billion	Includes a bidirectional SK Group arrangement covering NVIDIA's own memory purchases

Several major AI infrastructure counterparties — OpenAI, SoftBank, CoreWeave, Nebius, Oracle — have materially increased leverage since 2025.

Sources

- Stock Analysis, “NVIDIA (NVDA) Stock Price & Overview” and historical prices. <https://stockanalysis.com/stocks/nvda/>
- Fortune, “Markets experience new DeepSeek shock after MoonShot AI releases Kimi K3,” July 17, 2026.
- Bloomberg, “US Stocks Rebound From Selloff as Nvidia, Other Chipmakers Rise,” July 20, 2026.
- Bloomberg, “Moonshot’s Kimi K3 May Be More About Memory Than Compute,” July 20, 2026.
- Axios, “Exclusive: Nvidia’s Jensen Huang defends Chinese AI amid Kimi panic,” July 22, 2026.
- Axios, “Nvidia reignites ‘circular’ AI concerns as it weighs OpenAI financing guarantee,” July 27, 2026.

Prices are closing prices from public market data providers and may differ by a few cents across sources. Index moves and market-value estimates are as reported in the financial press.

The same tables are provided as a workbook, one worksheet per exhibit.

Copyright © 2026 Asaf Manela, Washington University in St. Louis, Olin Business School. All rights reserved. No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means without the permission of the copyright holder.